Participatory Censorship in Authoritarian Regimes

Tony Zirui Yang*

August 2, 2022

Abstract

Contrary to the conventional top-down view of government censorship, ordinary citizens in authoritarian regimes frequently participate in censorship by reporting online content. This study theorizes a bottom-up perspective of censorship in authoritarian regimes and analyzes its prevalence and consequences on public opinion toward censorship in the case of China. I argue that public participation increases support for censorship by diminishing the government's responsibility and strengthening citizens' perceived empowerment. Using an original survey in China, I show that more than half of the respondents have flagged content online and such participation is positively correlated with support for government censorship. I further conducted an experiment embedded in custom-engineered, simulated social media pages. Consistent with my theory, respondents that are encouraged to report simulated posts display significantly higher support for government censorship. My study highlights the role of ordinary citizens in facilitating authoritarian control and the normalization of repressive policies such as censorship.

Keywords: Censorship, Authoritarian Regime, China, Public Participation, Public Opinion

Word Count: 9641

*PhD Candidate, Department of Political Science, Washington University, St Louis, Campus Box 1063, One Brookings Drive, MO 63130, Email: yangzirui@wustl.edu

Introduction

Conventional wisdom of government censorship in authoritarian regimes regards it as a top-down tool of the authoritarian government to silence criticism and prevent collective actions that destabilize the regime (Gueorguiev and Malesky 2019; King, Pan and Roberts 2013; Miller 2018; Pop-Eleches and Way 2021; Roberts 2018; Shadmehr and Bernhardt 2015). In light of such understanding, existing research has extensively studied how unpopular government censorship can cause backlash (Pan and Siegel 2020; Roberts 2018, 2020), and how ordinary citizens circumvent and resist unpopular censorship activities (Chang et al. 2022; Chen and Yang 2019; Gläßel and Paula 2020; Han 2018; Hobbs and Roberts 2018; Roberts 2020).

Yet, contrary to conventional views, citizens in authoritarian regimes frequently participate in censorship by reporting online content. In Turkey, for example, Twitter users systematically weaponize the report function against political opponents (Tufekci 2017). In Russia, hundreds of users falsely report supporters of Ukraine on Facebook just before the Russian invasion of Ukraine (Nimmo and Agranovich 2022). In China, social media platforms explicitly ask their users to watch and report each other (Jiang 2021). Moreover, contrary to the conventional depiction of government censorship as unpopular, surveys across the world demonstrate significant levels of popular support for government censorship, especially in authoritarian regimes like China (Dickson 2016; Wang and Mark 2015), Russia (Nisbet, Kamenchuk and Dal 2017), and Middle Eastern monarchies (Martin, Martins and Wood 2016; Wike and Simmons 2015). *Why are many citizens in authoritarian regimes supportive of government censorship? Is it possible that participating in censorship increases citizens' support for it?*

This study takes a new step in this research field by exploring the consequences of a novel, bottom-up perspective of censorship. I theorize that when ordinary users actively participate in the censorship process by flagging online content that they disapprove of, their support for censorship increases. Specifically, public participation diminishes the government's respon-

sibility for censorship. When it is common for users to participate in censorship, citizens will be less likely to think that the government is responsible for censorship events. Even if citizens do not like certain censorship events, they might think that it is due to other users' reports, and are therefore less likely to blame the government. In addition, public participation increases citizens' perceived empowerment by censorship. When users can report posts that they disapprove of, they are less likely to regard themselves as victims of censorship. Instead, they are more likely to think censorship benefits them by limiting the information they do not like. As a result, censorship is less likely to be viewed as top-down political suppression imposed by the government upon ordinary users, and public opinion toward government censorship becomes more positive.

Using an original online survey in China, I provide the first descriptive analysis of the prevalence of public participation in the censorship process in authoritarian regimes. I find that more than half of the respondents self-report having flagged online content to be removed from the platforms. More than a quarter have reported online content at least once every few months. In particular, half of the "participating respondents" have the experience of reporting political content, including political news, opinions, and rumors. Such public participation in censorship is prevalent across several demographic groups and the political spectrum. In addition, I also find a significant and positive correlation between participation and support for censorship.

To causally test the theory, I conducted an original online experiment where participants are exposed to a custom-engineered, simulated social media environment. The simulated social media page is not interactive, meaning each respondent still completes the study independently. During the experiment on the simulated social media page, I give respondents in the treatment group an "encouragement" to participate in censorship by reporting the social media posts they encounter. I then use an instrumental variable analysis to estimate the local average treatment effect (LATE) of such experimentally induced censorship participation. Consistent with my theory, participation in censorship significantly increases individuals'

support for government censorship. I also find suggestive evidence that merely providing respondents with the option to report online content can increase support for censorship.

This study contributes to two streams of literature on authoritarian politics: government censorship and public participation. First, my study implies that censorship in authoritarian regimes should be perceived as a symbiotic relationship between the government and citizens. While the extant literature examines censorship in authoritarian regimes solely from a top-down angle (Gueorguiev and Malesky 2019; King, Pan and Roberts 2013; Lorentzen 2014; Miller 2018; Roberts 2018; Shadmehr and Bernhardt 2015), my study provides a novel theory to explain how this relationship could also be viewed from a bottom-up perspective. I show that ordinary users widely participate in the censorship process and such participation, in turn, shapes their opinion toward government censorship. More importantly, the bottom-up perspective helps to reconcile the empirical puzzle for why repressive apparatus such as censorship continues to garner widespread popular support in authoritarian regimes (Dickson 2016; Mitchell and Walker 2021; Wang and Mark 2015; Wike and Simmons 2015).

Second, I extend the literature on public participation in authoritarian regimes by showing a novel and perhaps insidious form of public participation. A wealth of literature has illustrated the causes and consequences of public participation in quasi-democratic institutions (Distelhorst and Hou 2017; Gandhi 2008; Gueorguiev 2021; He and Warren 2011; Manion 2015; Stromseth et al. 2017; Truex 2016, 2017), and contentious social movements (Fu and Distelhorst 2018). My study shows that public participation is also important in the implementation of repressive policies like censorship. Encouraging citizens' participation in the censorship process consolidates public support for authoritarian regimes' repressive apparatus. This semblance of public participation ironically suppresses individual rights to free speech and contributes to the durability of authoritarian regimes in the Internet era.

Participatory Censorship: A Bottom-Up Perspective

Although censorship in authoritarian regimes involves multiple levels of governments and non-governmental agencies (Han 2018; Lorentzen 2014; Tai and Fu 2020), it has always been understood as a top-down process imposed by the state and social media platforms upon ordinary citizens. Roberts (2018) defines censorship as "the restriction of the public expression of or public access to information *by authority* [emphasis added] when the information is thought to have the capacity to undermine the authority by making it accountable to the public." Han (2018) also regards censorship as "tools used *by the state* [emphasis added] to limit the boundaries of online expression." Even though recent research has found that the purposes of government censorship extend beyond silencing political opposition to rewarding regime supporters (Esberg 2020), there is nevertheless a consensus that censorship is a top-down political tool of authoritarian regimes. However, this traditional top-down view of censorship overlooks the important role ordinary users play in censorship. In this study, I explore a novel, bottom-up perspective of censorship where ordinary citizens participate in censorship by reporting online content that they disapprove of.

Many authoritarian regimes have long histories of public participation in repressive policies and political campaigns. During the Cultural Revolution of China, for instance, ordinary citizens reported their friends, colleagues, and even families, voluntarily or involuntarily, to the communist government as "counter-revolutionaries," which often led to brutal state repression of the reported persons (Dikötter 2016; Jiang 2021; Thurston 1984; Yang 2021). Similar reporting behaviors were also common in other dictatorships (Gregory 2009).

The emergence of new communication technologies has not diminished the prevalence and significance of public participation in authoritarian repression. Today, almost every social media platform has the function to report online content. Despite legitimate reasons to allow users to flag inappropriate and socially harmful content, governments and citizens can and do take advantage of such a feature to achieve their political goals (Nimmo and Agranovich 2022). For example, Tufekci (2017) documents that Turkish users on Twitter report political

opponents as "spam" in an organized manner. Such mass reporting of opponents is often successful in getting accounts suspended temporarily or even permanently (Tufekci 2017). In China, participation in censorship is especially prevalent. Official statistics claim that the government received over 163 million censorship requests in 2020 alone. Public figures ranging from outspoken dissidents like Fang Fang, to entertainment stars like Zhang Zhehan, to even government propagandist Hu Xijin, were all targets of these reports and many of their accounts were subsequently closed down (BBC News 2020). Some observers even highlight the disturbing parallels between the rising "report culture" online and the Cultural Revolution in the Mao era (BBC News 2020; Cook 2019; Jiang 2021).

The Chinese government welcomes such public input and even encourages citizens to report online content (Cook 2019; Jiang 2021). Since the establishment of the Cyberspace Administration of China (CAC) in 2014, encouraging public participation in censorship has been one of its main objectives. The CAC has an entire department, namely *jubao zhongxin*, dedicated to soliciting and handling censorship requests from ordinary users. It even ran an official propaganda campaign in Shanghai in 2020 to promote public participation in censorship (CNS 2020), and reward citizens for being "peer informants" (Cook 2019).

Why do authoritarian governments allow and even encourage such public participation? Theories of political participation in authoritarian regimes highlight the challenges of authoritarian governments in gathering various types of valuable information, such as public policy preference, popular grievance, and potential social unrest (Truex 2016). Due to the lack of democratic institutions, authoritarian governments rely on alternative quasi-democratic institutions (Distelhorst and Hou 2017; Gandhi 2008; Gueorguiev 2021; He and Warren 2011; Manion 2015; Stromseth et al. 2017; Truex 2016, 2017), and contentious social movements (Fu and Distelhorst 2018), to gauge public opinion and social unrest.

In the context of censorship, authoritarian governments need to identify messages that are threatening to the regime (Gueorguiev and Malesky 2019; King, Pan and Roberts 2013). Yet, even in China, one of the most sophisticated censorship regimes, it is still difficult for

the government to completely control the Internet (Roberts 2018). To effectively censor from a large amount of information, the Chinese government allegedly hires millions of censors (King, Pan and Roberts 2017), uses automated keyword filtering (Han 2018; Ng 2015), and uses a "friction" strategy to make all sensitive content harder to get (Roberts 2018).

Encouraging ordinary users to participate in the censorship process is another strategy that authoritarian governments can and do employ to alleviate the information gathering problem on the Internet (Nimmo and Agranovich 2022; Tufekci 2017). Instead of monitoring every corner of the entire Internet, users' reports provide valuable signals for the government and social media firms to conduct censorship. Sina Weibo, for example, establishes an algorithm (*quanzhong*) that takes users' reports into account when determining the publicity and censorship of every account and every post (Cook 2019; Jiang 2021). By mobilizing millions of ordinary users to participate in the censorship process and establishing an algorithm that automatically filters widely reported posts, the cost of monitoring the Internet is significantly reduced. As a result, the dynamic of censorship in China is no longer solely top-down control but a mixture of top-down control and bottom-up participation.

Participation and Public Support for Censorship

How does public participation in censorship affect public opinion toward it? Existing research has shown that government censorship of political opposition is likely to cause public backlash against the regime (Pan and Siegel 2020; Roberts 2018, 2020). However, such a backlash effect is based on two assumptions: citizens believe (1) that the government is responsible for the censorship event, and (2) that censorship negatively affects their welfare. I argue that public participation in the censorship process diminishes the government's responsibility. In addition, citizens are more likely to feel empowered by participating in censorship and are less likely to view themselves as victims of censorship. As a result, public support for censorship is more likely to increase.

Diminishing Government's Responsibility

Authoritarian governments are cautious not to take the blame for suppressing political opposition. The Chinese government, in particular, regularly outsources repression to non-state actors to avoid the impression that it is directly repressing political dissent (Han 2018; Ong 2022). For example, Ong (2022) finds that the Chinese state hires various kinds of violent thugs and non-violent brokers, as well as mobilizes the mass to carry out everyday repression and achieve social control. By outsourcing everyday repression, the regime plausibly denies any wrongdoings and evades political accountability.

In the case of censorship, existing studies have identified two main "scapegoats" for the Chinese government. First, the Chinese government strategically mobilizes its pro-regime base to fabricate millions of posts and counter online critics while avoiding direct censorship of political opposition (Chen and Xu 2017; Han 2015; King, Pan and Roberts 2017; Miller 2018). Such a strategy diverts the focus of regime critics from the government to their pro-regime counterparts. Second, the Chinese government also outsources censorship to social media platforms (Han 2018), and openly criticizes them when backlash against online censorship occurs (Miller 2018). As a result, the blame for censorship is diverted from the government to social media companies.

Public participation in the censorship process further contributes to the diffusion of the responsibility of the government. First, public participation gives the government a passive image in censorship events. It creates the perception that the government is not the initiator of censorship and is merely responding to public demand (Luo and Li forthcoming). As such, even if citizens do not like certain censorship events, they might think that it is due to other users' reports, and therefore less likely to blame the government. Second, public participation also makes censorship activities seem more popular and normal. Traditionally, censorship is often perceived as government suppression of free speech against the will of the people (Pan and Siegel 2020). When the public is actively participating in the censorship process, it creates the perception that censorship is not just the will of the government but

also the will of many ordinary users, and it is normal for online content to be removed (Yang 2022). This perception of popularity and public acceptance itself can reduce citizens' inclination to speak out against the censorship apparatus (Huang and Cruz 2021).

Increasing Citizens' Perceived Empowerment

In addition to diffusing the responsibility, authoritarian governments also frequently engage in justification of their repression to increase citizens' perceived empowerment by repression. The goal of such justification is to persuade the public that they are not the victims, but rather the beneficiaries of repressive policies (Esberg 2018; Williamson and Malik 2019). The Chinese government also tries to shift public perception of censorship in this regard. For example, Yang (2022) finds that the Chinese government expands the range of censorship targets beyond politically threatening content to other seemingly harmless non-political content such that the public no longer regards censorship as political repression but rather a normal government policy that benefits ordinary users.

Because participation in censorship allows ordinary users to report content that they disapprove of, it reduces citizens' perception that they are victims of censorship by increasing their perceived control over the censorship outcome. Because ordinary users can report online content that affects subsequent censorship activities, it delegates some, albeit small, amount of censorship power to ordinary users and therefore increases their perception of control over the content being censored. Existing studies in social and political psychology have found that individuals with a low perception of control are more susceptible to conspiracies, less politically engaged, and more likely to attribute influence or blame to political enemies (Nyhan and Zeitzoff 2018; Smith 2021; Sullivan, Landau and Rothschild 2010; Whitson and Galinsky 2008). When censorship is solely a top-down process imposed upon ordinary users, they are more likely to have cynical views of the censorship apparatus, perceive themselves as victims of censorship events, and blame the government for such negative consequences. On the contrary, because of the increased perception of control over censorship outcomes,

individuals are more likely to view government censorship activities as enforcing their censorship preferences. Hence, they are more likely to feel empowered by censorship and view themselves as beneficiaries of censorship rather than victims.

The changes in the perceived government responsibility and citizen empowerment, in turn, likely affect how ordinary citizens view censorship. From the perspective of ordinary users, their participation in the censorship process redefines the government's role as an arbitrator of conflicting interests on the Internet rather than a manipulator of public opinion. They are less likely to view censorship as top-down political suppression imposed upon them, and they are more likely to treat censorship as a tool they can use to their advantage: to suppress political opponents (Luo and Li forthcoming). As a result, ordinary citizens are more likely to support the censorship apparatus.

Hypothesis: As individuals participate more in the censorship process, they should display greater levels of support toward government censorship.

Study 1: Online Survey

Although participatory censorship occurs in various types of autocracies (Tufekci 2017), I focus on China, one of the most sophisticated censorship regimes, to illustrate the features and consequences of participatory censorship. To gauge the prevalence of public participation in censorship among the Chinese public and its correlation with support for censorship, I conducted an original online survey in December 2021.[1] The survey recruited 1,124 respondents via a Chinese online survey platform. After they agreed to participate, the respondents were re-directed to an American-based website, Qualtrics, where they completed the survey anonymously. I employ a quota sampling strategy to recruit respondents with a wide range of socioeconomic backgrounds. However, like many other online surveys in China, the sample will inevitably be younger and better-educated than the general Internet population (Huang

[1]This study was approved by the Institutional Review Board (IRB) at the researcher's home institution and was pre-registered on Open Science Framework.

2018; Pan and Xu 2020). To address this concern, I weight the survey sample such that the weighted sample resembles the Chinese Internet census in terms of gender, rural/urban location, region, and age distributions.[2]

Measurement

Participation in Censorship

To measure *Participation in Censorship*, I directly ask the respondents if they have reported online content and speeches before. Respondents are given five choices: never, once or twice only, once per few months, once per month, and multiple times per month. The responses are coded on a five-point scale. Although surveys in authoritarian regimes often face social desirability bias problems, it is unlikely to be a concern here. As discussed in previous sections, censorship participation is a prevalent behavior in China and the Chinese government holds a favorable view toward such public participation. Hence, it is unlikely that respondents are fearful of reporting their previous behaviors.

For those respondents who self-report having participated before, I further query the specific types of content they requested to censor by providing a list of content for them to choose from. Specifically, I am interested in whether they have requested censorship of political content, including political news, political commentary, political opinions, political rumors, and foreign media coverage of China. I also include other non-political content such as entertainment, advertisement, vulgar language, and pornography in the list.

Outcome Variables

The main dependent variable of the analysis is *Support for Censorship*. I ask the respondents if they agree that the government should actively control the Internet and remove content that it deems inappropriate. I further measure their *Support for Censorship of Political Content* and *Support for Censorship of Non-Political Content* by asking whether the government should control online discussion of government policies and party leadership, as

[2]For more information on the survey sample and weighting, see appendix.

well as entertainment stars and popular culture.

Control Variables

I include two different sets of control variables: demographic covariates and predisposition covariates. Demographic covariates include *education*, *age*, *gender*, and *urban/rural location*. Predisposition covariates include *party membership*, *political interests*, *political ideology*, and *economic ideology*. The first two covariates are commonly used in surveys in China (Huang 2018). The third and fourth questions measuring ideology are directly borrowed from Pan and Xu (2020).

Results: Prevalence of Public Participation

How prevalent is public participation in censorship among the Chinese public? Figure 1 presents the distribution of self-report participation in censorship using the weighted sample. As shown in the first bar on the left, more than 50% of the weighted sample self-report having previously participated in censorship. More than 25% have participated at least once every few months. These results demonstrate that public participation in censorship is prevalent among Chinese internet users. Such behavior is especially ubiquitous among the younger generations who are regularly exposed to online discussions about the use of reporting as a strategy to censor other opposing opinions (Luo and Li forthcoming). Almost three-quarters of respondents under 20-year-old self-report having such censorship experience, and around two-thirds of respondents in their 20s report similar experiences. For these young people on the Internet, flagging online content is both common and normal.

Moreover, the participation rate does not vary significantly by political ideology.[3] To be clear, there is not a clear political cleavage in the Chinese society Pan and Xu (2020), but existing studies have found groups that self-identify as more liberal and critical of the regime compared to other groups that are more supportive of the regime. These ideologically

[3]I measure political ideology by asking whether they agree or disagree that: From a long-term perspective, multi-party democracy is suitable for China. The question is borrowed from Pan and Xu (2020). Those who agree with the statement are categorized as liberal, whereas those who disagree are conservative.

diverging groups both interact with each other as well as the state censorship apparatus (Han 2018). The results here show that even politically liberal citizens choose to weaponize the "report" function for political gains. This implies that even though many individuals might object to the authoritarian system as a whole, they do not resist participating in the repressive apparatus so long as it serves their immediate interests.

Figure 1: Distribution of Self-Report Participation in Censorship

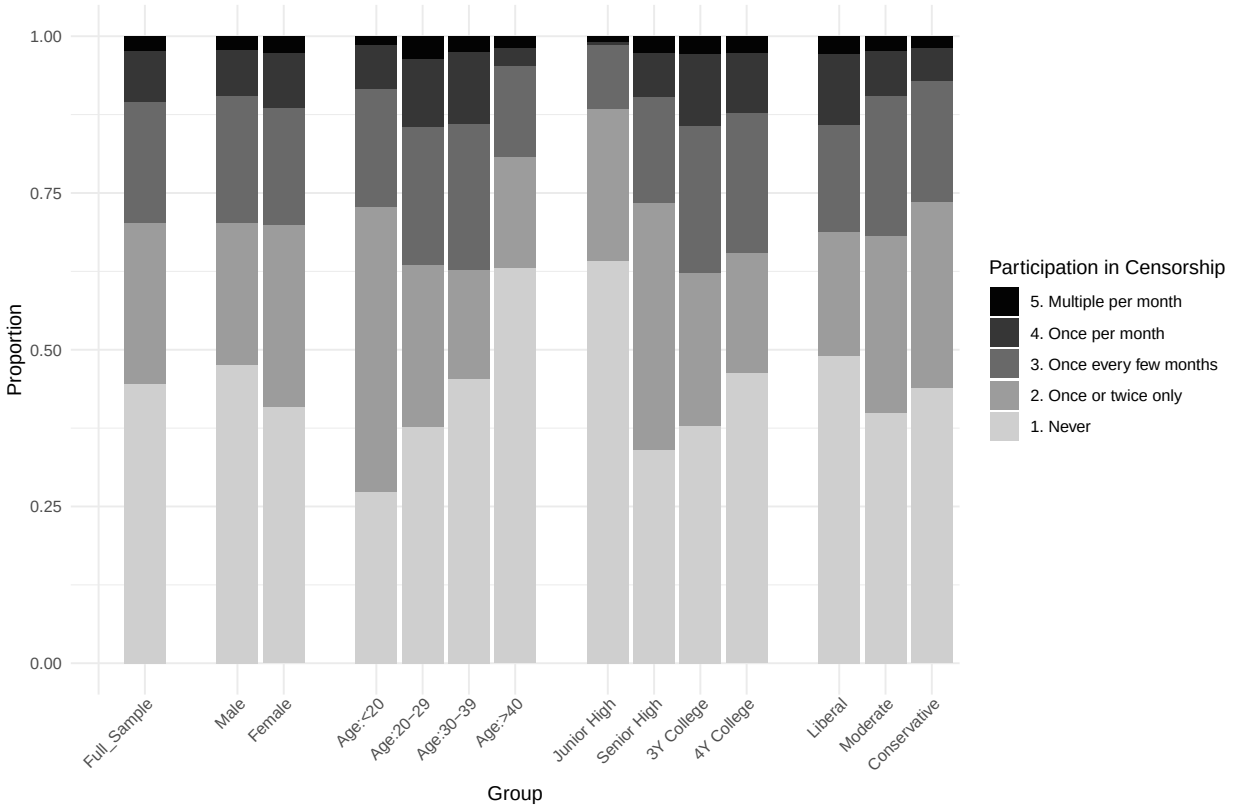

Note: All observations are weighted by gender, rural/urban location, region, and age group. Unweighted sample shows a slightly higher proportion of respondents self-reporting participation in censorship.

What specific content did these "participating respondents" report? Most "participating respondents" (around 90%) have flagged inappropriate content online including pornography, advertisement, and vulgar language. A significant proportion has also reported political content. In general, around 50% of the "participating respondents," or 25% of all respondents, self-report having flagged political content. Similar to the general trend described above,

younger generations and the better-educated are significantly more likely to participate in the censorship of political content. Both liberals and conservatives have engaged in political censorship with a slightly higher participation rate among liberals. About one-third of the "participating respondents" have reported entertainment and cultural content. In sum, the descriptive analysis demonstrates significant levels of public participation in censorship overall and across different demographic groups and different categories of online content.

Participation and Censorship Support

Do individuals with higher levels of participation in censorship hold more favorable views toward government censorship? Figure 2 reports the results of regressing *Support for Censorship* on *Participation in Censorship* using OLS models. Because the proportion of respondents who participate in censorship multiple times per month is too small (around 3% of the sample), I combine the top two highest participation categories. Thus, the independent variable is measured on a four-point scale. Consistent with the hypothesis, censorship participation positively and significantly correlates with individuals' support for censorship.[4] This suggests that individuals who have flagged more content in the past are more likely to believe that the government should actively remove content it deems harmful or inappropriate.

Such a correlation also holds for censorship of specific categories of online content. Respondents who have participated more in censorship before are more likely to believe that the government should actively control political news and discussions, as well as entertainment content and discussions of popular culture.

Taken together, the results from the observational analyses support my claim that public participation in censorship is common among the Chinese public. About one-half of the respondents self-report participated in censorship. Moreover, such participation is positively

[4]Controlling for all covariates, the magnitude of the effect of censorship participation is equivalent to 8.05% of a standard deviation of the dependent variable. Using the original five-point scale to measure censorship participation does not significantly change the results. As a robustness check, I also re-coded the variable as a binary variable indicating whether the respondent has reported before. The results are consistent with the main analyses as well.

Figure 2: Correlation between Participation in Censorship and Support for Censorship

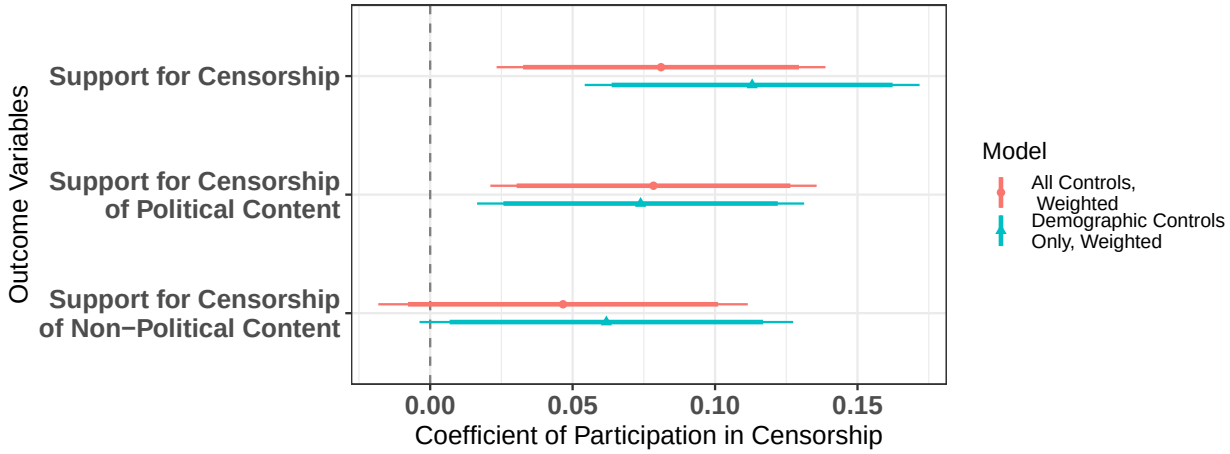

Note: All outcome variables are measured on a five-point Likert scale. Participation in censorship is initially measured on a five-point scale: never participated, once or twice only, once per few months, once per month, and multiple times per month. But the top two categories are combined in the analyses due to a low proportion of respondents reporting multiple times per month. Bars indicate 90% and 95% confidence intervals.

correlated with support for censorship agreeing with my main hypothesis. These findings highlight the relevance of the bottom-up perspective of censorship in authoritarian regimes. If one-half of the online population is actively engaged in reporting online content, it is important to analyze its consequences in order to fully understand how citizens in authoritarian regimes view censorship and repressive government policies in general.

However, the analyses above do not rule out the possibility of a reversed causal arrow: individuals who support censorship in the first place are more likely to participate in it. To provide stronger support for the effect of public participation on support for censorship, I will need to randomly assign participation in censorship and causally test my hypothesis.

Study 2: Survey Experiment

Building on the first study, I conducted an original online survey experiment on a custom-engineered, simulated social media page.[5] The simulated social media environment is not

[5]This study was approved by the Institutional Review Board (IRB) at the researcher's home institution and was pre-registered on Open Science Framework.

interactive, meaning each participant still completed the survey independently. The goal of the experiment is to test the effect of participation in censorship on public support for government censorship. Because it is difficult to directly manipulate censorship behavior, I use an instrumental variable approach, i.e., intention-to-treat design. I give respondents in the treatment group an encouragement to participate in censorship and measure the local average treatment effect (LATE) on support for censorship among respondents who complied with my encouragement (Aronow and Carnegie 2013; Marbach and Hangartner 2020). Such an approach allows me to estimate the causal effects of censorship participation induced by the encouragement treatment. In addition, I also test the intention-to-treat effect of the encouragement treatment by comparing the group means of the two groups.

Conducting such an experiment in a simulated social media environment has several advantages. First, it avoids ethical concerns associated with a similar field experiment. The current political climate in China is hostile toward political field experiments and might put both participants and researchers at higher risk of authoritarian repression. In addition, encouraging respondents to participate in censorship in the real world will be normatively undesirable and might further contribute to the reporting culture on the Chinese Internet. Conducting the experiment in a simulated setting limits the potential negative impact of the research. Finally, reporting behaviors are usually not publicly observable in the real world and, as confirmed in my experiment, a substantial proportion of respondents object to the idea of reporting and refuse to participate in censorship even after encouragement. An experiment embedded in a simulated environment enables me to measure whether respondents participated, estimate the treatment effects among compliers, and identify those who refuse to participate regardless of treatment.

Procedure

The survey experiment was conducted in June 2022 in China. Similar to the online survey in study 1, I recruited 4,008 respondents from a Chinese online survey platform and then

directed them to Qualtrics, where they completed the survey anonymously. As was the case with study 1, the sample covers a wide range of socioeconomic backgrounds but is younger and better-educated than the general Internet population.

The experiment consists of three parts. First, I measure pre-treatment covariates. Second, I randomly assign respondents to one of the three groups: a control group, a treatment group, and a blank control group. All respondents read the same ten social media posts on a heatedly debated current event, Xuzhou chained woman incident,[6] on my simulated social media page. All posts are adapted from real Weibo posts with modified user names and avatars. Respondents are fully informed about how these posts are created from real posts, and therefore no deception is used. Among the ten posts, five of them could be seen as pro-government or nationalistic, while the other five could be seen as anti-government or pro-individual rights. The order of the posts is randomized. After they read the social media posts, I measure respondents' support for censorship.

On the simulated social media page, I build multiple buttons that the respondents can click under each post. In both the control group and the treatment group, these buttons are: "Like," "Share," "Comment," and "Report." To manipulate respondents' participation in censorship in the simulated social media environment, I give the respondents in the treatment group an "encouragement" to use the "Report" button. Specifically, respondents in the treatment group were shown the following paragraph:

> *We are especially interested in what posts you want to report. Please choose at least two posts that you think should be removed by the Internet regulator, and press the Report button to let us know.*

In addition to the control and the treatment groups, I also include an additional blank control group in which I only build three buttons that the respondents can click: "Like,"

[6]In January 2022, a video of a trafficked woman held in chains in a hut in Fengxian County, Xuzhou City for years went viral. The government and officials were heavily criticized for both causing such tragedy and trying to cover it up. However, the incident coincided with the 2022 Beijing Winter Olympics. Such timing prompted many patriotic regime supporters to argue that this is a conspiracy to defame China.

"Share," and "Comment." In other words, respondents in the blank control group do not have the opportunity to report. Naturally, they do not receive the "encouragement" to report as well. Comparing the blank control group to the control group tests the effect of providing the opportunity to participate in censorship on support for censorship. While this is not a direct test of the main hypothesis, it is a useful complement to the encouragement design and provides a stronger test of my theory. Detecting significant effects of merely providing the opportunity to participate might imply that the effect sizes would be substantially larger if respondents actually participated in the censorship process. Table 1 summarizes the treatment assignment.

Table 1: Treatment Summary

Groups	Blank Control	Control	Treatment
Buttons Under Simulated Posts	Like, Share, Comment	Like, Share, Comment, **Report**	Like, Share, Comment, **Report**
Encouragement Message	No	No	Yes

Notes: The instrumental variable analysis only includes the control and treatment groups. To test the effect of providing participation opportunities, I only include the blank control and the control groups, but not the treatment group.

Measurement

To measure respondents' participation, I use a binary variable to indicate whether they have clicked any of the "Report" buttons. Similar to the observational survey, my main dependent variable of the analysis is *Support for Censorship*. I also measure their *Support for Censorship of Political Content* and *Support for Censorship of Non-political Content*.

I use ten covariates to check the balance across the three experimental groups. They are also included in the regression analyses. Among the ten covariates, four are demographic variables, including *education*, *age*, *gender*, and *region*. These demographic covariates are widely used in experiments across different contexts. The remaining six covariates are pre-

disposition covariates including *party membership*, *nationalism*, *political interests*, *ideology*, *social media usage*, and *foreign connection*. All covariates are balanced across the three experimental groups.

Overall Results

I first report the overall results of the experiment by comparing the group means of the outcome variables across the three groups, and then elaborate on the instrumental variable analysis that more precisely identifies the causal effects of participation on support for censorship. As shown in Figure 3, I find a consistent increase in support for censorship first from the blank control group to the control group, and then from the control group to the treatment group.

Figure 3: Group Means of the Dependent Variables

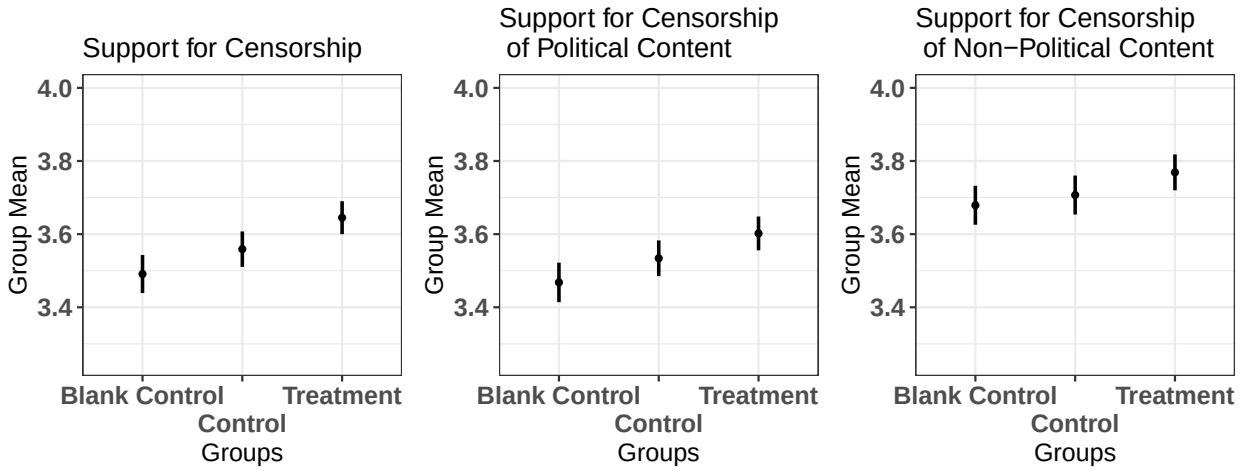

Note: All three outcome items are measured on a five-point Likert scale. In the left and center panels, the differences between treatment and control groups are significant at the 0.05 level; the differences between blank control and control groups are significant at the 0.1 level. In the right panel, the difference between the control and treatment groups is significant at 0.1 level; the difference between blank control and control groups is not significant.

Comparing the blank control group and the control group tests the effect of providing the opportunities to report online content on support for censorship. The only difference between the two groups, and therefore the "treatment," is that the control group has an additional

"Report" button under each post, whereas the blank control group does not. The additional "Report" button primes the respondents that ordinary citizens have the opportunity to participate in censorship. Similar to the logic laid out in the theory section, providing the opportunities to report could diminish perceived government responsibility for censorship and increase citizens' perceived empowerment by the censorship apparatus. Hence, I should expect a positive effect on individuals' support for censorship.

As shown in Figure 3, I find respondents who have the option to report simulated posts express higher support for censorship in general ($\beta = 0.067$, $p = 0.065$) and censorship of political content in particular ($\beta = 0.066$, $p = 0.076$). Both results are significant at the 0.1 level. I do not find a significant difference in support for censorship of non-political content, potentially because the topic I selected is a political one. Taken together, the results provide suggestive evidence that merely providing the opportunity to flag online content can increase individuals' support for censorship.

Next, I compare the difference between the control and the treatment groups. While both groups have the option to report, respondents in the treatment group received an encouragement to use the "Report" bottom. As I will describe in more detail in the instrumental variable analysis, the encouragement treatment increases the participation rate by around 20%. The difference in means between the treatment and the control groups represents the intention-to-treat effect of the encouragement treatment. Consistent with the main hypothesis, I find that respondents who are encouraged to flag online content display significantly higher support for government censorship in general ($\beta = 0.086$, $p = 0.011$), and support for censorship of political content in specific ($\beta = 0.068$, $p = 0.048$). The treatment effects on support for censorship of non-political content ($\beta = 0.062$, $p = 0.092$) is significant at the 0.1 level.

All in all, the overall results of the experiment support my central argument that higher levels of participation lead to higher support for the censorship apparatus. Both the institution that allows public participation in censorship and the encouragement that directly

increases reporting behaviors can generate significant support for the censorship apparatus. One caveat, however, is that neither of the treatments above is reporting behavior per se. Thus, to more precisely estimate the effect of reporting behaviors on censorship support, the next section introduces the instrumental variable analysis and reports its results.

Instrumental Variable Analysis

To directly identify the effect of participation in censorship, I use the encouragement treatment as an instrument to estimate the local average treatment effect (LATE) of participation in censorship (clicking the report button) on support for censorship. Formally:

$$\begin{aligned} \text{Clicking the Report Button}_i &= \alpha + \gamma \cdot \text{Encouragement Treatment} + \lambda Z_i + \epsilon_i \\ Y_i &= \xi + \beta \cdot \widehat{\text{Clicking the Report Button}_i} + \delta Z_i + \mu_i \end{aligned}$$

where Y_i is the outcome measure; Z_i is a vector of pre-treatment covariates; and β is the LATE (or the average treatment effect among compliers). This analysis only includes the treatment and control groups, but not the blank control group.

In the control group, 43% of the respondents (i.e. always-takers) clicked the "Report" buttons on the simulated social media page, whereas 64% of the respondents in the treatment group clicked. Hence, there are roughly 20% of the respondents (i.e. compliers) clicked the "Report" buttons due to the encouragement treatment. The instrumental variable analysis estimates the causal effects of censorship participation induced by the encouragement treatment (i.e. average treatment effect among compliers).

Table 2 reports the results from instrumental variable analyses and the LATE of participating in censorship in the simulated social media environment. Consider, first, column 2. After controlling for pre-treatment covariates, participation induced by the encouragement treatment significantly increases respondents' general support for government censorship ($\beta = 0.402$, $p = 0.017$). This again provides direct support for the main hypothesis and more importantly, it alleviates the concerns in study 1 that the causal arrow might be re-

versed. The magnitude of the treatment effect on support for censorship is equivalent to 46.3% of a standard deviation. It is a substantial increase considering the baseline support for censorship is already high in the control group (3.59 out of 5).

Censorship participation induced by the encouragement treatment also increases specific support for censorship of political content ($\beta = 0.305$, $p = 0.072$) and non-political content ($\beta = 0.319$, $p = 0.081$), although both coefficients are only significant at the 0.1 level. This might suggest that while in general, the results are consistent with the theoretical expectations, it is more difficult to move individuals' opinions on specific censored content.

Table 2: Local Average Treatment Effect (ATE among Compliers) of Participating in Censorship on Support for Censorship

	Support for Censorship		Support for Censorship of Political Content		Support for Censorship of Non-Political Content	
	(1)	(2)	(3)	(4)	(5)	(6)
Report Click	0.417**	0.402**	0.330*	0.305*	0.299*	0.319*
	(0.169)	(0.168)	(0.169)	(0.170)	(0.180)	(0.183)
Constant	3.379***	1.977***	3.392***	2.097***	3.579***	2.216***
	(0.092)	(0.193)	(0.092)	(0.195)	(0.098)	(0.211)
Covariates		✓		✓		✓
N	2,647	2,493	2,653	2,499	2,645	2,492

Notes: Dependent variables are indicated in column headings. Standard errors in parentheses. Report click is a binary variable indicating whether the respondents have clicked any of the "Report" buttons on the simulated social media page. All individual survey items were measured on a five-point scale.
*p < .1; **p < .05; ***p < .01

To check the robustness of the treatment effect, I use an alternative measurement of participation in censorship and re-run the instrumental variable analyses (see Online Appendix). Instead of a binary variable indicating whether the respondents clicked any of the "Report" buttons, I use the number of times the respondents clicked a "Report" button. Consistent with the main analyses, additional clicking of the "Report" buttons induced by the encouragement treatment significantly increases support for censorship.

Profiling Compliers

Although the results from the instrumental variable analyses show positive and significant treatment effects of censorship participation on support for government censorship, they only capture the effects of the treatment among a subgroup of the sample: the compliers. Hence, it is important to compare the background attributes of compliers and non-compliers (Marbach and Hangartner 2020).

As mentioned earlier, compliers only account for around 20% of the full sample, whereas never-takers and always-takers consist of about 36% and 43% respectively. Such a low proportion of compliers might potentially be a problem for the generalizability of the treatment effects. However, a comparison of pre-treatment covariates among compliers, never-takers, and always-takers suggests that it might be less concerning. Always-takers are more likely to be young and social media savvy, whereas never-takers are more likely to be old, nationalists with few foreign connections and social media exposure. Hence, those never-takers are likely to be more pro-regime and they are not necessarily the subgroup that the regime tries to win over by promoting participation in censorship. Compliers, in comparison, are young, better-educated, and less nationalistic people with more foreign connections. Hence, because of the similarity in background attributes between compliers and always-takers, it is reasonable to assume that the treatment effect of participation in censorship is likely to apply to always-takers. More importantly, compliers and always-takers are more likely to be on the Internet and therefore more likely to be the target population of my theory on public participation in censorship.

It is noteworthy, however, that the composition of compliers, always-takers, and never-takers might be affected by the selected topic and posts. In other words, some individuals who participate in censorship in the real world might not necessarily click the "Report" buttons in this specific experiment and vice versa. Thus, the profiling results above should be interpreted with caution and future research should test the same argument and hypothesis in different scenarios to improve generalizability.

Discussion

To sum up the findings in the experiment, both the institutional feature that allows public participation in censorship and the message that directly encourages participation increase support for censorship to various degrees. The instrumental variable analysis further demonstrates that the treatment effects among respondents who complied with the encouragement treatment are large and significant.

The implications of the experimental results are three-fold. First, it echoes the findings in the observational survey that public participation in the censorship process is indeed a prevalent phenomenon in authoritarian China. Always-takers account for more than 40% of the full survey sample, meaning a substantial proportion of the respondents voluntarily participate in the censorship process even without experimental treatment. Putting compliers and always-takers together, a simple encouragement message can lead to almost two-thirds of the respondents participating in censorship.

Second, the results reaffirm the positive relationship between widespread participation in the censorship process and public support for government censorship. It indicates that the button-up perspective of censorship has significant consequences on public opinion. A more comprehensive approach to authoritarian censorship that combines top-down control and bottom-up participation is needed to better understand the dynamic of censorship and authoritarian control. Moreover, the results suggest that while the government might benefit from merely providing participation access to the public, it is the encouragement to participate that eventually leads to a significant increase in censorship support.

Third, the profiling compliers analysis suggests that the effect is primarily driven by young and well-educated individuals with foreign connections. These individuals have traditionally been understood as more open to Western values and less susceptible to authoritarian controls (Huang 2015). The results of this experiment suggest that public participation in repressive policies like censorship might be an effective mechanism through which the authoritarian regime generates support for its repressive apparatus from these young, well-

educated, and global-minded citizens, who are traditionally less inclined to do so. By encouraging censorship participation among these subgroups, the regime successfully sustains popular support for its repressive apparatus across generations.

Alternative Mechanisms

A limitation of both studies, however, is that I do not find strong support for either of the proposed mechanisms that public participation increases support for censorship by diminishing the government's responsibility and strengthening citizens' perceived empowerment. In both studies, I measure respondents' *Perceived Empowerment* by asking whether ordinary people are the victims or the beneficiaries of the current censorship apparatus; and respondents' *Perceived Government Responsibility* by asking whom they think should be responsible for the censored content: the netizens, the government, or the platforms. To test the mechanisms in study 1, I perform OLS regressions to test the correlations between *Participation in Censorship* and the two mechanism variables. In study 2, I run the same instrument variable models as in the main analysis to identify the LATE of the encouragement treatment on the two mechanism variables.

As shown in Table 3, the estimates of all four models are in the correct direction as my theory predicts: respondents who participated in censorship are less likely to believe that the government is responsible for censorship events and more likely to believe that they are empowered by censorship. Yet, none of the estimates is statistically significant. One possible explanation for the insignificant results is the lack of variation in responses. The variances of both mechanism outcomes are lower than censorship support questions and responses are concentrated on the middle choices.

Given that the support for the proposed mechanisms is weak, what alternative mechanisms might explain the main findings? One possible alternative mechanism is cognitive dissonance (Festinger 1957). It claims that encountering inconsistencies between one's be-

Table 3: The Effects of Participation in Censorship on Mechanism Variables

	Study 1		*Study 2*	
Dependent Variable	Government Responsibility	Perceived Empowerment	Government Responsibility	Perceived Empowerment
Model	OLS	OLS	IV	IV
	(1)	(2)	(3)	(4)
Participation	−0.074	0.029		
	(0.049)	(0.061)		
Report Click			−0.282	0.211
			(0.349)	(0.395)
Constant	3.603***	4.126***	4.387***	3.360***
	(0.283)	(0.357)	(0.403)	(0.461)
Covariates	✓	✓	✓	✓
Weighted Sample	✓	✓		
N	1,074	1,047	2,507	2,291

Notes: Dependent variables are indicated in column headings and measured on a five-point Likert scale. Standard errors in parentheses.
Participation in censorship is initially measured on a five-point scale, but the top two categories are combined in the analyses. Report click is a binary variable indicating whether the respondents have clicked any of the "Report" buttons on the simulated social media page.
*p < .1; **p < .05; ***p < .01

liefs and behaviors causes psychological discomfort, which motivates individuals to reduce the inconsistencies. While individuals might initially oppose censorship, they might still participate in censorship because they want to remove content they do not like. Such behavior is inconsistent with their belief and therefore creates cognitive dissonance. As such, individuals are motivated to change their attitude toward government censorship in order to reduce psychological discomfort.

Another possible alternative mechanism is the system justification theory (Jost 2020). It states that people are motivated to defend and justify the social, economic, and political arrangements on which they depend (Jost 2020). This implies potential psychological needs for individuals in authoritarian regimes to justify initially unpleasant behaviors and experiences they encountered. Hence, when individuals are constantly engaged in reporting online content, they will unconsciously justify such behavior and the censorship apparatus.

Although I cannot directly test these alternative mechanisms, one observable implication from both the cognitive dissonance theory and the system justification theory is that participation in censorship should have a stronger effect on increasing individuals' support for the censorship apparatus than increasing support for the authoritarian regime more broadly. In study 2, I asked three regime support questions: *Regime Satisfaction*, *Regime Assessment*, and *Regime Trust*.[7] As shown in Figure 4, there is no observable difference between the control group and the blank control group across the three regime support items. The differences between the control and the treatment groups are also insignificant at the conventional level. In short, the results do show a stronger effect of participation on censorship support than regime support, and provide suggestive evidence for the alternative mechanisms. However, directly testing the cognitive dissonance theory and the system justification theory remains a task for future research.

Figure 4: Group Means of Regime Support Survey Items

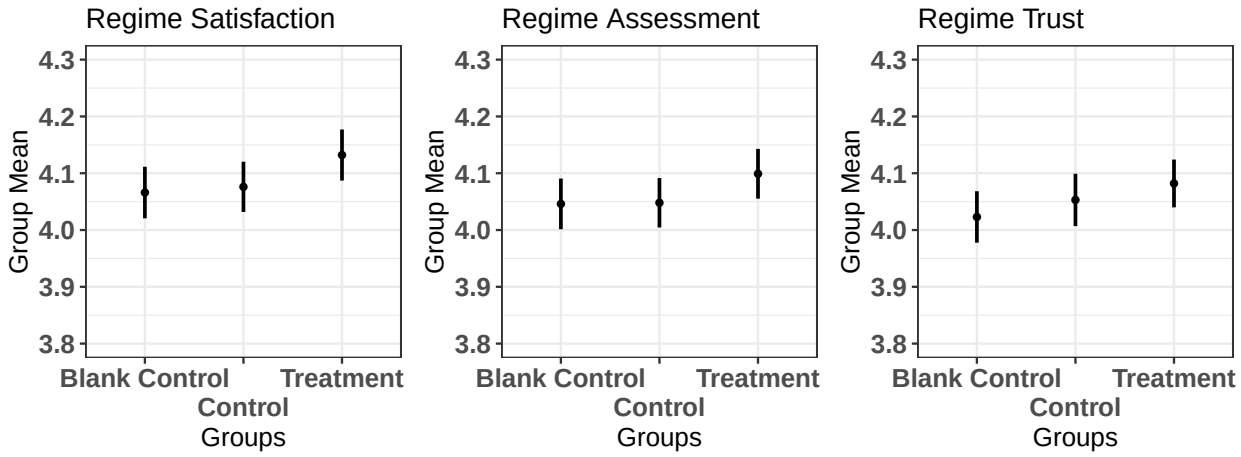

Note: All three outcome items are measured on a five-point Likert scale. The differences between the blank control and the control groups are insignificant across the three items. The difference between the control and treatment groups is only significant at the 0.1 level for regime satisfaction, but not significant for the other two items.

[7]For more information about the regime support survey items, see the Online Appendix.

Conclusion

Censorship in authoritarian regimes has long been understood as top-down suppression of political opposition. This study provides a novel theorization, empirical description, and quantitative analysis of bottom-up public participation in censorship in authoritarian regimes. Using two original online surveys in China, one observational and one experimental, I demonstrate that public participation in censorship is both prevalent and significant in shaping public opinion toward the censorship apparatus.

The findings in this study highlight the discrepancy between the common understanding of repressive authoritarian apparatus, such as censorship, in the Western world and how ordinary citizens in authoritarian regimes perceive and interact with these repressive apparatus. For many citizens in China, censorship and other repressive institutions have been normalized as part of the rules of political life (Yang 2022). Therefore, instead of fighting against the rules, they take advantage of it and use censorship to suppress opposing views (Luo and Li forthcoming; Tufekci 2017) . Such a mentality has significant downstream political implications as citizens no longer view the regime as the oppressor, but rather as a powerful arbitrator of censorship demands and whom citizens should win over in their internal fights against fellow citizens.

Beyond autocracies, censorship has become an important social issue in many democracies including the United States, and public participation in censorship has also become more prevalent. Twitter, for example, introduced a community-based bottom-up content moderation project called "Birdwatch." Although democracies might care about different policy implications, such as electoral integrity, compared with their authoritarian counterparts, it is still important to analyze the consequences of these content moderation projects involving public participation, because the balancing act of fighting misinformation and preserving freedom of speech is difficult yet critical for sustaining a healthy democracy.

References

Aronow, Peter M and Allison Carnegie. 2013. "Beyond LATE: Estimation of the Average Treatment Effect with an Instrumental Variable." *Political Analysis* 21(4):492–506.

BBC News. 2020. "China where Fang Fang is From: the Report Culture is Spreading along with the Rise of Nationalism." *BBC News Chinese Version* .
URL: *https://www.bbc.com/zhongwen/simp/chinese-news-52560990*

Chang, Keng-Chi, William R. Hobbs, Margaret E. Roberts and Zachary C. Steinert-Threlkeld. 2022. "COVID-19 Increased Censorship Circumvention and Access to Sensitive Topics in China." *Proceedings of the National Academy of Sciences* 119(4).

Chen, Jidong and Yiqing Xu. 2017. "Why do Authoritarian Regimes Allow Citizens to Voice Opinions Publicly?" *The Journal of Politics* 79(3):792–803.

Chen, Yuyu and David Y Yang. 2019. "The Impact of Media Censorship: 1984 or Brave New World?" *American Economic Review* 109(6):2294–2332.

CNS. 2020. "2020 Shanghai shi wangluo jubao xuanchuan yue zhengshi qidong [2020 Shanghai Internet Reporting Promotion Month Officially Launched]." *China News Agency* .
URL: *https://www.chinanews.com/sh/2020/09-02/9280670.shtml*

Cook, Sarah. 2019. "Analysis: How the Chinese Communist Party Is Incentivizing Repression." *China Media Bulletin* 133.

Dickson, Bruce. 2016. *The Dictator's Dilemma: The Chinese Communist Party's Strategy for Survival*. Oxford University Press.

Dikötter, Frank. 2016. *The Cultural Revolution: A People's History, 1962—1976*. Bloomsbury Publishing USA.

Distelhorst, Greg and Yue Hou. 2017. "Constituency Service Under Nondemocratic Rule: Evidence from China." *The Journal of Politics* 79(3):1024–1040.

Esberg, Jane. 2018. "The Audience of Repression: Killings and Disappearances in Pinochet's Chile." *Available at SSRN 3246120* .

Esberg, Jane. 2020. "Censorship as Reward: Evidence from Pop Culture Censorship in

Chile." *American Political Science Review* 114(3):821–836.

Festinger, Leon. 1957. *A Theory of Cognitive Dissonance*. Vol. 2 Stanford university press.

Fu, Diana and Greg Distelhorst. 2018. "Grassroots Participation and Repression Under Hu Jintao and Xi Jinping." *The China Journal* 79(1):100–122.

Gandhi, Jennifer. 2008. *Political Institutions Under Dictatorship*. Cambridge University Press.

Gläßel, Christian and Katrin Paula. 2020. "Sometimes Less is More: Censorship, News Falsification, and Disapproval in 1989 East Germany." *American Journal of Political Science* 64(3):682–698.

Gregory, Paul R. 2009. *Terror by Quota*. Yale University Press.

Gueorguiev, Dimitar. 2021. *Retrofitting Leninism: Participation Without Democracy in China*. Oxford University Press.

Gueorguiev, Dimitar D and Edmund J Malesky. 2019. "Consultation and Selective Censorship in China." *The Journal of Politics* 81(4):1539–1545.

Han, Rongbin. 2015. "Defending the authoritarian regime online: China's "voluntary fifty-cent army"." *The China Quarterly* 224:1006–1025.

Han, Rongbin. 2018. *Contesting Cyberspace in China: Online Expression and Authoritarian Resilience*. Columbia University Press.

He, Baogang and Mark E Warren. 2011. "Authoritarian Deliberation: The Deliberative Turn in Chinese Political Development." *Perspectives on Politics* 9(2):269–289.

Hobbs, William R and Margaret E Roberts. 2018. "How Sudden Censorship Can Increase Access to Information." *American Political Science Review* 112(3):621–636.

Huang, Haifeng. 2015. "International Knowledge and Domestic Evaluations in a Changing Society: The Case of China." *American Political Science Review* 109(3):613–634.

Huang, Haifeng. 2018. "The Pathology of Hard Propaganda." *The Journal of Politics* 80(3):1034–1038.

Huang, Haifeng and Nicholas Cruz. 2021. "Propaganda, Presumed Influence, and Collective

Protest." *Political Behavior* pp. 1–24.

Jiang, Jue. 2021. "The Eyes and Ears of the Authoritarian Regime: Mass Reporting in China." *Journal of Contemporary Asia* 51(5):828–847.

Jost, John T. 2020. *A Theory of System Justification*. Harvard University Press.

King, Gary, Jennifer Pan and Margaret E. Roberts. 2013. "How Censorship in China Allows Government Criticism but Silences Collective Expression." *American Political Science Review* 107:326–343.

King, Gary, Jennifer Pan and Margaret E Roberts. 2017. "How the Chinese government fabricates social media posts for strategic distraction, not engaged argument." *American political science review* 111(3):484–501.

Lorentzen, Peter. 2014. "China's Strategic Censorship." *American Journal of Political Science* 58(2):402–414.

Luo, Zhifan and Muyang Li. forthcoming. "Participatory Censorship: How Online Fandom Community Facilitates Authoritarian Rule." *New Media & Society* .

Manion, Melanie. 2015. *Information for Autocrats: Representation in Chinese Local Congresses*. Cambridge University Press.

Marbach, Moritz and Dominik Hangartner. 2020. "Profiling Compliers and Noncompliers for Instrumental-Variable Analysis." *Political Analysis* 28(3):435–444.

Martin, Justin D, Ralph J Martins and Robb Wood. 2016. "Desire for Cultural Preservation as a Predictor of Support for Entertainment Media Censorship in Saudi Arabia, Qatar, and the United Arab Emirates." *International Journal of Communication* 10:23.

Miller, Blake. 2018. Delegated Dictatorship: Examining the State and Market Forces behind Information Control in China PhD thesis.

Mitchell, Amy and Mason Walker. 2021. "More Americans Now Say Government Should Take Steps to Restrict False Information Online than in 2018." *Pew Research Center* 18.

Ng, Jason Q. 2015. "Politics, Rumors, and Ambiguity: Tracking Censorship on WeChat's Public Accounts Platform." *Munk School of Global Affairs* .

Nimmo, Ben and David Agranovich. 2022. "Meta's Adversarial Threat Report, First Quarter 2022.".

Nisbet, Erik C, Olga Kamenchuk and Aysenur Dal. 2017. "A Psychological Firewall? Risk Perceptions and Public Support for Online Censorship in Russia." *Social Science Quarterly* 98(3):958–975.

Nyhan, Brendan and Thomas Zeitzoff. 2018. "Fighting the past: Perceptions of control, historical misperceptions, and corrective information in the Israeli-Palestinian conflict." *Political Psychology* 39(3):611–631.

Ong, Lynette H. 2022. *Outsourcing Repression: Everyday State Power in Contemporary China*. Oxford University Press.

Pan, Jennifer and Alexandra A. Siegel. 2020. "How Saudi Crackdowns Fail to Silence Online Dissent." *American Political Science Review* 114:109–125.

Pan, Jennifer and Yiqing Xu. 2020. "Gauging Preference Stability and Ideological Constraint under Authoritarian Rule." *Working Paper, Available at SSRN* .

Pop-Eleches, Grigore and Lucan A Way. 2021. "Censorship and the Impact of Repression on Dissent." *American Journal of Political Science* .

Roberts, Margaret E. 2018. *Censored: Distraction and Diversion Inside China's Great Firewall*. Princeton University Press.

Roberts, Margaret E. 2020. "Resilience to Online Censorship." *Annual Review of Political Science* 23:22.1–22.19.

Shadmehr, Mehdi and Dan Bernhardt. 2015. "State Censorship." *American Economic Journal: Microeconomics* 7(2):280–307.

Smith, Brianna A. 2021. "It's All Under Control: Threat, Perceived Control, and Political Engagement." *Political Psychology* .

Stromseth, Jonathan R, Edmund J Malesky, Dimitar D Gueorguiev, Lai Hairong and Carl Brinton. 2017. *China's Governance Puzzle: Enabling Transparency and Participation in a Single-Party State*. Cambridge University Press.

Sullivan, Daniel, Mark J Landau and Zachary K Rothschild. 2010. "An existential function of enemyship: Evidence that people attribute influence to personal and political enemies to compensate for threats to control." *Journal of personality and social psychology* 98(3):434.

Tai, Yun and King-wa Fu. 2020. "Specificity, Conflict, and Focal Point: A Systematic Investigation into Social Media Censorship in China." *Journal of Communication* 70(6):842–867.

Thurston, Anne F. 1984. "Victims of China's Cultural Revolution: The Invisible Wounds: Part I." *Pacific Affairs* 57(4):599–620.

Truex, Rory. 2016. *Making Autocracy Work: Representation and Responsiveness in Modern China*. Cambridge University Press.

Truex, Rory. 2017. "Consultative authoritarianism and its limits." *Comparative political studies* 50(3):329–361.

Tufekci, Zeynep. 2017. *Twitter and Tear Gas*. Yale University Press.

Wang, Dakuo and Gloria Mark. 2015. "Internet Censorship in China: Examining User Awareness and Attitudes." *ACM Transactions on Computer-Human Interaction (TOCHI)* 22(6):1–22.

Whitson, Jennifer A and Adam D Galinsky. 2008. "Lacking control increases illusory pattern perception." *Science* 322(5898):115–117.

Wike, Richard and Katie Simmons. 2015. "Global Support for Principle of Free Expression, But Opposition to Some Forms of Speech." *Pew Research Center* 18.

Williamson, Scott and Mashail Malik. 2019. "Contesting narratives of repression: Experimental evidence from Sisi's Egypt." *Journal of Peace Research* .

Yang, Jisheng. 2021. *The World Turned Upside Down: A History of the Chinese Cultural Revolution*. Farrar, Straus and Giroux.

Yang, Tony Zirui. 2022. "Normalization of Censorship: Evidence from China." *Working Paper, Available at SSRN 3835217* .